
Priority-Based Scheduling

- Jobs have priorities
- Scheduler selects highest-priority ready job
- Two possibilities:
 - **Dynamic-priority scheduling**
 - Different jobs of task may be assigned different priorities
 - Example
 - Job $J_{i,k}$ of task T_i has higher priority than job $J_{j,m}$ of T_j
 - Job $J_{i,l}$ of T_i has lower priority than job $J_{j,n}$ of T_j
 - Static-priority scheduling

Dynamic-Priority Scheduling


- Consider a task's characteristics:
 - ❑ Phase
 - ❑ Period
 - ❑ Execution time
 - ❑ Relative deadline
- What derived characteristics does each job have?
 - ❑ Release time
 - ❑ Absolute deadline
- Main goal: schedule jobs so that no job misses deadline
 - ❑ What would be the most straightforward way to do this?

Earliest Deadline First (EDF)

- Whenever scheduling decision needs to be made
 - Job with earliest absolute deadline given highest priority
- When must scheduling decisions be made?
 - 1) When currently executing job completes
 - 2) When new job is released/activated
- Consequence of 2)
 - Jobs can be interrupted during execution
 - I.e., jobs can be preempted

Example

- Derive preemptive EDF schedule for task set below
 - $T_1(8, 2)$, $T_2(2, 5, 1, 5)$, $T_3(10, 2)$ and $T_4(6, 2)$



Earliest Deadline First (EDF)

Theorem 4-1: [Liu and Layland] When preemption is allowed and jobs do not contend for resources, the EDF algorithm can produce a feasible schedule of a set **J** of jobs with arbitrary release times and deadlines on a processor if and only if **J** has feasible schedules.

In other words, preemptive EDF is **optimal**


Notes:

- Applies even if tasks are not periodic
- If periodic, relative deadline could be $<$, $=$ or $>$ period
- Precedence constraints are allowed

Proof sketch for Theorem 4-1

Any feasible schedule of \mathbf{J} can be systematically transformed into an EDF schedule

Suppose parts of two jobs J_i and J_k are executed out of EDF order:



This situation can be corrected by performing a “swap”:




Proof sketch for Theorem 4-1

Inductively repeating this can eliminate all out-of-order violations

Resulting schedule may still have idle intervals even when job is ready:




Such idle intervals can be eliminated by moving some jobs forward:




What about non-preemptive EDF?

Consider a system of three jobs J_1 , J_2 , and J_3 such that
 $(r_1, e_1, d_1) = (0, 3, 10)$, $(r_2, e_2, d_2) = (2, 6, 14)$, $(r_3, e_3, d_3) = (4, 4, 12)$.

Here's a feasible schedule:



Now, what happens if we use non-preemptive EDF?



So, what's the conclusion?

Theorem: Non-preemptive EDF is not optimal.

Question: Does this mean preemptive EDF is always better than non-preemptive EDF in practice?

Not necessarily!

EDF optimality proof assumes there is no penalty due to preemption

Note: from now on, “EDF” means preemptive EDF, unless specified otherwise

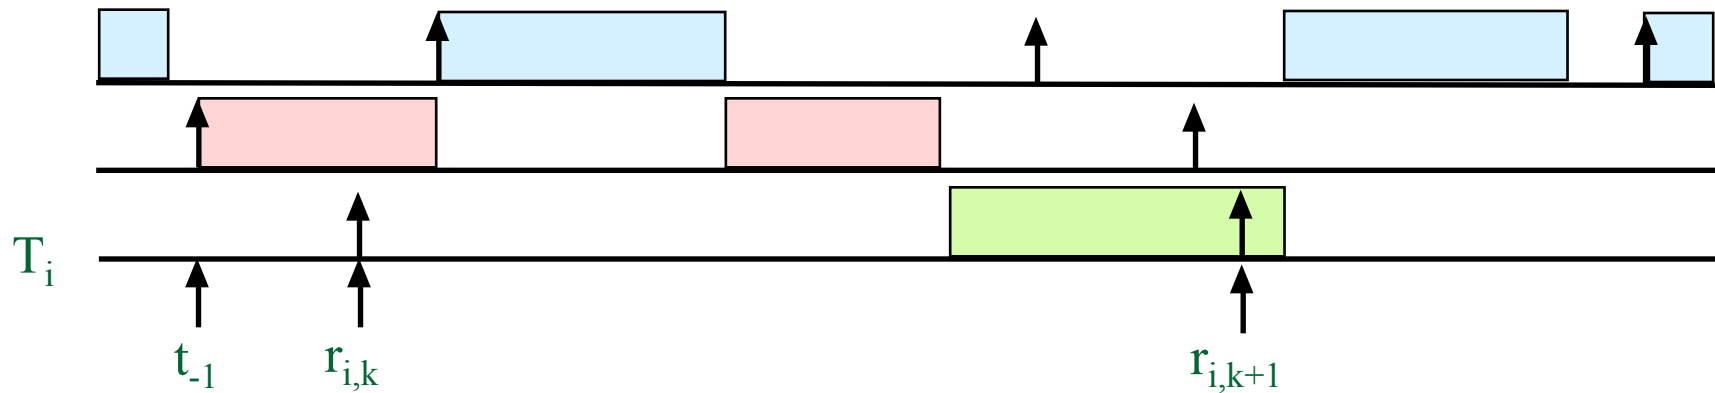
Schedulability test for (preemptive) EDF

- Obvious observation
 - Total utilization (U) must be less than or equal to 1 (100%)
- But, if $U \leq 1$, is the task set surely schedulable?
 - The contrapositive of a true statement is always true
 - Statement: If $U \leq 1$, then the task set is schedulable
 - Contrapositive: if task set is not schedulable, then $U > 1$
 - Our approach
 - Start with assumption that task set is not schedulable
 - Work backwards and prove the contrapositive

Proving the contrapositive

Assume task set is not schedulable

Let $J_{i,k}$ be the first job to miss its deadline.



this is the last “idle instant” (each task whose next job has a deadline at or before $r_{i,k+1}$ either has no ready job, or has *just* released a job, and no job with a deadline after $r_{i,k+1}$ executes in $[t_{-1}, r_{i,k+1}]$.)

Proving the contrapositive

$J_{i,k}$ missed its deadline...

→ demand placed on processor in $[t_{-1}, r_{i,k+1})$ by jobs with deadlines $\leq r_{i,k+1}$ is greater than the available processor time in $[t_{-1}, r_{i,k+1}]$

$r_{i,k+1} - t_{-1}$ = available processor time within $[t_{-1}, r_{i,k+1}]$

$r_{i,k+1} - t_{-1} <$ demand placed on processor within $[t_{-1}, r_{i,k+1})$
by jobs with deadlines $\leq r_{i,k+1}$

$r_{i,k+1} - t_{-1} < \sum_{j=1}^N \{ \# \text{ jobs of } T_j \text{ with deadlines } \leq r_{i,k+1} \text{ released within } [t_{-1}, r_{i,k+1}) \} e_j$

$r_{i,k+1} - t_{-1} < \sum_{j=1}^N \left\lfloor \frac{r_{i,k+1} - t_{-1}}{p_j} \right\rfloor e_j \Rightarrow r_{i,k+1} - t_{-1} < \sum_{j=1}^N \frac{r_{i,k+1} - t_{-1}}{p_j} e_j$

Proving the contrapositive

We have

$$r_{i,k+1} - t - 1 < \sum_{j=1}^N \frac{r_{i,k+1} - t - 1}{p_j} e_j$$

Canceling $r_{i,k+1} - t - 1$ yields

$$1 < \sum_{j=1}^N \frac{e_j}{p_j}$$

i.e.,

$$1 < U$$

Note: This proof is valid even if **relative deadlines are larger than periods**

So, what's the conclusion?

If $U > 1$, task set is not schedulable

If $U \leq 1$, task set is schedulable

Here's the formal theorem:

Theorem 6-1: [Liu and Layland] A system T of independent, preemptable, periodic tasks with relative deadlines equal to their periods can be feasibly scheduled (under EDF) on one processor if and only if its total utilization U is at most one.

EDF with Deadlines < Periods

If deadlines are less than periods then $U \leq 1$ is no longer a sufficient schedulability condition

Consider two tasks such that, for both, $e_i = 1$ and $p_i = 2$

If both have deadlines at 1.9, then the system is not schedulable, even though $U = 1$.

Here, densities used instead of utilizations

Definition: The **density of task T_k** is defined as $\delta_k = \frac{e_k}{\min(D_k, p_k)}$

The **density of the system** is defined as $\Delta = \sum_{k=1..N} \delta_k$

EDF with Deadlines < Periods

Theorem 6-2: A system **T** of independent, preemptable, periodic tasks can be feasibly scheduled on one processor if its density is at most one.

Proof is similar to that for Theorem 6-1

Note: This theorem only gives **sufficient** condition

We refer to the following as the **schedulability condition for EDF:**

$$\sum_{k=1}^n \frac{e_k}{\min(D_k, p_k)} \leq 1$$

Proof of non-tightness


Note that $\Delta > 1$ doesn't imply non-schedulability

Example:

Consider two tasks $T_1 = (2, 0.6, 1)$ and $T_2 = (5, 2.3)$

$$\Delta = 0.6/1 + 2.3/5 = 1.06$$

Nonetheless, we can schedule this task set under EDF:




Properties of EDF

- Priority calculation is dynamic at task level
 - Priorities of different jobs of task are different
 - Priority of specific job does not change during its execution
- Priority calculation can also be job level dynamic...

Least Laxity First (LLF)

- **Definition:** At any time t , the slack (or laxity) of a job with deadline d is equal to $d - t$ minus the time required to complete the remaining portion of the job



- **LLF Scheduling:** job with smallest laxity has highest priority

Optimality of LLF

Theorem 4-3: When preemption is allowed and jobs do not contend for resources, the LLF algorithm can produce a feasible schedule of a set \mathbf{J} of jobs with arbitrary release times and deadlines on a processor if and only if \mathbf{J} has feasible schedules.

Proof is similar to that for EDF

Question: Which of EDF and LLF would be preferable in practice?